

CPT S 483 Project 2: Conway's Game of Life

You are allowed to work as individuals or in teams of size two for this project. Make only one submission for each team project. The other person makes an “empty” submission with just a word document that is the cover sheet (from the website) indicating what your team is.

In this project you will implement the *Conway's Game of Life* in MPI. The game is a cellular automaton containing $m*n$ cells arranged as m rows and n columns. Each cell can be in one of the two states at any given point of time – *alive* or *dead*. In the initialization step, we assign the state for each cell by a toss of a coin (i.e., random with equal probability). After that, game simply keeps progressing from one generation to another without needing any further user input.

During the k^{th} generation, all cells determine their respective states based on the following simple **rules**:

- i) A living cell that has less than 2 living neighbors dies (as if caused by underpopulation, or loneliness, or simply boredom!);
- ii) A living cell with more than 3 living neighbors also dies (as if caused by overcrowding);
- iii) A living cell with either 2 or 3 living neighbors lives on to the next generation;
- iv) A dead cell with 3 living neighbors will be “born” (or “reborn” if it was ever alive before) and become alive.

Therefore, the decision to live or die obviously depends on the states of a cell's neighbors and its own current state. As a rule, we will use the states from *the preceding* generation to make the decisions for the current generation. This implies that the states of cells in the current generation can be updated independent of one another and in any arbitrary order without affecting the output.

Neighborhood: Since the automaton is structured as a grid (or mesh), each internal cell will contain exactly 8 neighbors: {N,S,E,W,NE,NW,SE,SW}. For cells that are on the boundary of the automaton (first row/column, last row/column), neighborhood will include wrapped around neighbors from the other end of the matrix – for instance, the first row will treat the last row as its preceding row, while the last column will treat the first column as its preceding column. The corner cells will use the corresponding cell in the diagonally opposite corner as substitute for their corresponding missing diagonal neighbor entries.

Your task for this assignment is to write an efficient implementation for the Conway's Game of Life in C/C++/MPI and analyze its scalability (both analytically and empirically). For the purpose of this assignment, we will make the following assumptions:

- that the automaton's matrix is always a **square** with n rows and n columns;
- that $n \gg p$ and n is divisible by p (the number of processes);
- that the number of generations to simulate is specified by the user at input. Let this parameter be denoted by G ;

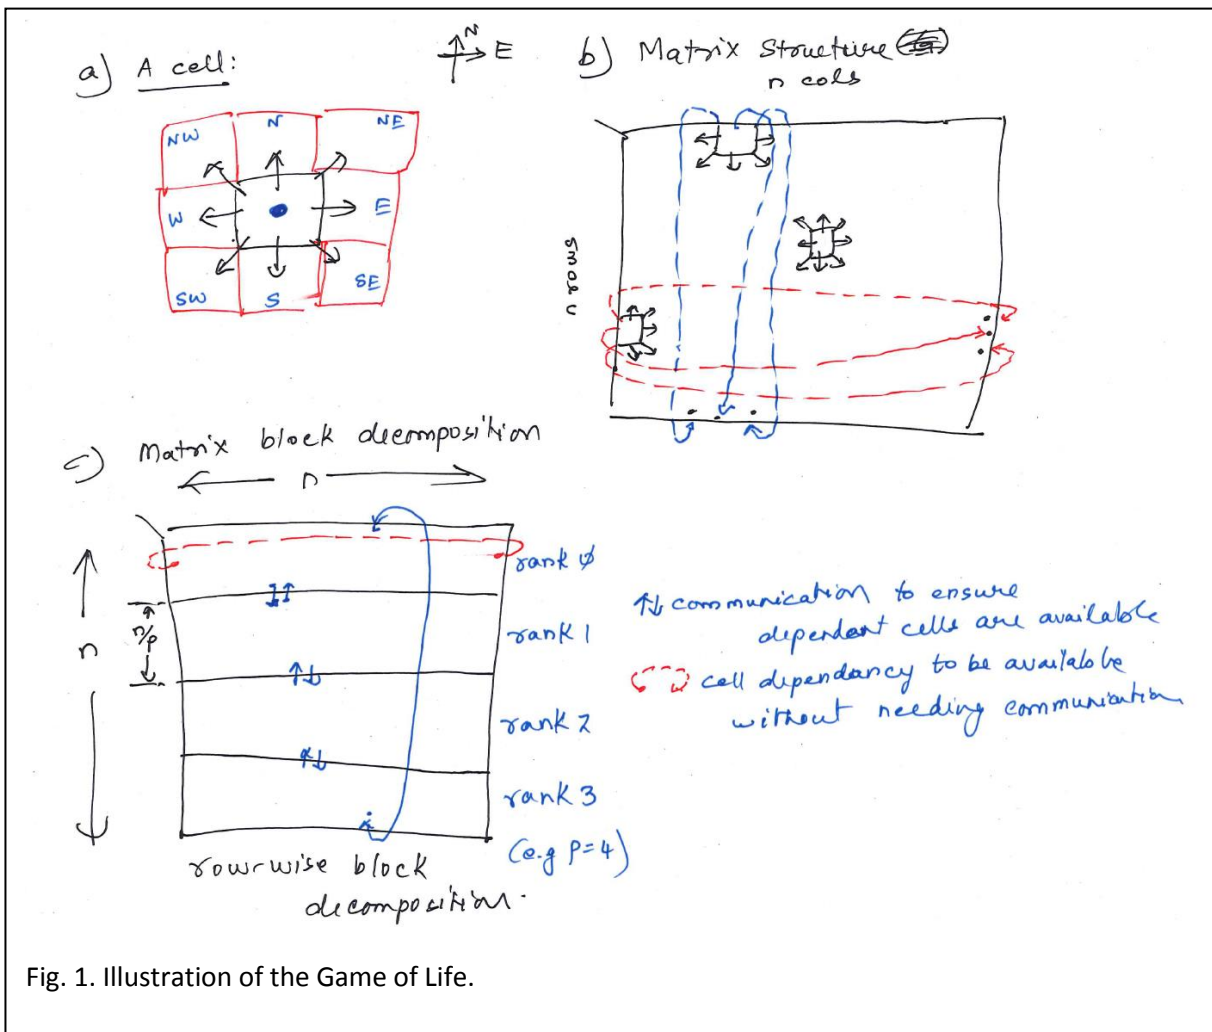


Fig. 1. Illustration of the Game of Life.

Parallel implementation:

There are multiple ways in which the Game of Life can be parallelized. One way is to decompose the input matrix into roughly equal sized smaller, nonoverlapping squares such that each process owns a distinct, part of the matrix. Another way is to decompose the matrix into disjoint batches of rows (or columns) such that each process gets n/p rows (or columns). The latter approach, which is typically referred to as *block decomposition*, is better suited for implementation in distributed memory (why?). So we will use that approach in the implementation. Note that in both approaches, each rank will get n^2/p cells.

The pseudocode for your algorithm is as follows:

Input: User specifies parameters n, G

Initialization:

- *GenerateInitialGoL()*¹: This function should generate the initial matrix in parallel so that by the end of the call, the entire matrix is generated and stored in a distributed manner (with rank i owning rows $[i*(n/p) \dots (i+1)*(n/p)-1]$). To do this, there are two steps:
 - Rank 0 first generates p random numbers (in the interval of 1 and BIGPRIME²) and **distributes** them such that the i^{th} random number is handed over to rank i .
 - Next, using the assigned random number as the “seed”, each rank generates a distinct sequence of n^2/p random values, again in the same interval. Each generated random number is used in the following manner to fill the initial matrix: if the k^{th} random number is even, then the k^{th} cell being filled in the local portion of the matrix is marked with status=Alive; otherwise its state=Dead. (This is one way of generating the matrix randomly. If you prefer to do it other ways, that’s fine too but make sure you use different random seeds in different processes to avoid the problem of generating matrix replicas across processes.) Also note that if you initialize this way, the initially filled matrix is not necessarily compliant with the rules of the GoL. That’s okay. You will fix it in the first iteration of your simulation. Read on...
- *Simulate()*: This function actually runs the Game of Life. The simulation is done for G generations (i.e., iterations). Within each iteration, rank i determines the new states for all the cell that it owns. To determine the new state of a cell, write a function called *DetermineState()* that takes a cell coordinate and returns its new state (alive or dead) based on the GoL rules described above. For this to happen, however, you need to first make sure all cell states have access to their neighboring entries – i.e., do the necessary communication to satisfy these dependencies *a priori*. You also need to ensure that all processes are executing the same generation at any given time. This can be done using a simple *MPI_Barrier* at the start of every generation.
- *DisplayGoL()*: As the simulation proceeds it is desirable to display the contents of the entire matrix to visualize the evolution of the cellular automaton. However, doing this after every step of simulation could be cumbersome for large n . Therefore, we will do this infrequently – i.e., display once after every x number of generations. I will let you decide the value of x based on the timings you see for a specific input size. To do the display, write a function called *DisplayGoL()* which first aggregates (i.e., gathers) the entire matrix in rank 0 (“root”) and then displays it. (Note, displaying from every separate rank independently is possible but could possibly shuffle up the prints out of order.) As an alternative to each process sending its entire local matrix to the root you could choose to send only the alive entries. This approach would provide some communication savings if the matrix is sparse with very few alive entries. However for simplicity if you want to just send the whole local matrix that is OK for this assignment.

Reporting

¹ Note that I have just mentioned the function names here. You are free to decide on the arguments as necessary in your implementation.

² Some reasonably big prime number – e.g. 93563, 68111, etc.

- Timing: Your code should measure the following times for reporting purposes:
 - Total runtime (excluding the time for display) to run the simulation for the user-specified G generations
 - Average time per generation (excluding time for display)
 - Average time per the display function
 - Time for the different communication steps (i.e., if you are using say *MPI_Barrier*, *MPI_Send*, *MPI_Recv*, *MPI_Gather*, *MPI_Scatter*, *MPI_Bcast*, *MPI_Reduce*, etc. in different places of your code, then attach a timer specific to each of those calls and report them separately at the end – i.e., one timer per MPI function).

These timings should be reported (from each rank) at the end before the program terminates.

Experiments

To test the scalability of your code, devise the following experiments: Measure the average time per generation (excluding display time) as a function of the number of processes, for varying input sizes but a fixed number of generations (G). Think of a table for this purpose, with rows for different input sizes ($n \times n$) and columns for different number of processes. To change the input size, grow n in powers of 2, starting from an input that is as small as 4×4 and going as large as $2^{10} \times 2^{10}$ or more as dictated by your runtime. Vary the number of processes in powers of 2: e.g., $\{1, 2, 4, 8, 16\}$ in our glx cluster.

Using the above table, generate three plots:

- i) **Speedup** in the Average runtime per generation (Y-axis) vs. Number of processes (X-axis) – with different curves for different input sizes ($n \times n$);
- ii) **Efficiency** curves for the above chart;
- iii) Plot the breakdown (in %) of the average time per generation to indicate how much of it was spent in “Computation” (i.e., the time to update the local matrix and anything else that is not included any communication) vs. “Communication” (i.e., the sum of the times spent by each rank in communicating primitives). Show this breakdown for varying number of processes for a fixed input size.

Interpretation: Briefly state your observations about your results – do they meet your analytical expectations? If not, why not? Do you see ways to optimize this further?

Deliverables (zipped into one zip file with your name on it):

- i) Full source code
- ii) Report in PDF (preferred) or Word that shows all scaling results and your interpretation of those results.